


STATS 315B Final



Suppose cost function $C(w) = \sum_{i=1}^N c_i(w)$ where $c_i(w) = \frac{1}{2} [y_i - F(x_i)]^2$. Further, let node outputs for layer k be $\vec{z}_k = (z_{1k}, \dots, z_{mk})$ with input node $\vec{x} = (x_1, \dots, x_n)$. Finally, let the edge from z_{ik} to $z_{j(k+1)}$ have weight $w_{ij}^{(k)}$.

Then, using sigmoid function $S(x) = \frac{1}{1+e^{-x}}$

$$z_{j1} = x_j \quad \forall j \in [1, n]$$

$$z_{0k} = 1 \quad \forall k \in [1, L]$$

$$z_{m(k+1)} = S\left(w_{0m}^{(k)} + \sum_{i=1}^{m_k} w_{im}^{(k)} z_{ik}\right) \quad \forall k \in [1, L-1]$$

$$z_1^{L+1} = w_{01}^{(L)} + \sum_{i=1}^{m_L} w_{i1}^{(L)} z_{iL}$$

Then, for each node z_{mk} , error e_{mk} is

$$e_{1(L+1)} = F(x_i) - y_i \quad (m=1)$$

$$e_{mk} = z_{mk}(1-z_{mk}) \left(\sum_{i=1}^{m_k} \omega_{mj}^{(k)} e_{i(k)} \right) \quad (m \neq 1)$$

The e_{mk} for $m \neq 1$ is computed recursively starting from $k=L$ and decrementing by 1 till $k=1$ is reached.

Now, we propose

$$[1] \quad e_{mk} = z_{mk}(1-z_{mk}) \frac{\partial C_i}{\partial z_{mk}}$$

~~We~~ We prove this via induction:

Base Case: $k=L$,

$$\begin{aligned} e_{mL} &= \omega_{m1}^{(L)} z_{mL}(1-z_{mL}) e_{1(L+1)} \\ &= \omega_{m1}^{(L)} z_{mL}(1-z_{mL}) (F(x_i) - y_i) \\ &= z_{mL}(1-z_{mL}) \frac{\partial C_i}{\partial z_{mL}} \quad \checkmark \end{aligned}$$

Inductive step: Before proceeding, we use the Sigmoid function property $S'(x) = S(x)(1-S(x))$, which is easily verified. So now, assuming [1] holds for $k+1$, we show it holds for k .

$$e_{mk} = z_{mk} (1 - z_{mk}) \left(\sum_{i=1}^{m_{k+1}} w_{mi}^{(k)} e_{i(k+1)} \right)$$

$$= z_{mk} (1 - z_{mk}) \left(\sum_{j=1}^{m_{k+1}} w_{mj}^{(k)} S'(\omega_{oj}^{(k+1)}) + \sum_{j=1}^{m_k} w_{ij}^{(k)} z_{jk} \right) \frac{\partial C_i}{\partial z_{jk(k+1)}}$$

Using a forementioned property:

$$= z_{mk} (1 - z_{mk}) \left(\sum_{j=1}^{m_{k+1}} \frac{\partial z_{jk(k+1)}}{\partial z_{mk}} \cdot \frac{\partial C_i}{\partial z_{jk(k+1)}} \right)$$

$$= z_{mk} (1 - z_{mk}) \frac{\partial C_i}{\partial z_{mk}}$$

∴ we are done

So to compute the gradient $\nabla(\vec{w})$, we have

$$\frac{\partial C_i}{\partial w_{rs}^{(k)}} = z_{rk} S'(\omega_{os}^{(k)} + \sum_{j=1}^{m_k} w_{js}^{(k)} z_{jk})$$

$$= z_{rk} S'(\omega_{os}^{(k)} + \sum_{j=1}^{m_k} w_{js}^{(k)} z_{jk}) \frac{\partial C_i}{\partial z_{s(k+1)}}$$

$$= z_{rk} z_{s(k+1)} (1 - z_{s(k+1)}) \frac{\partial C_i}{\partial z_{s(k+1)}}$$

$$= z_{rk} e_{s(k+1)}$$

So, the algorithm works as follows:

1.) Set ~~e_i~~ $e_i^{(k+1)} = F(x_i) - y_i$

2.) For $k = 1$ to K , $k++$,

- Compute $\frac{\partial \mathcal{L}_i}{\partial \omega_{rs}^{(k)}} = z_{rk} e_i^{(k+1)}$
for all edges.

- For $m = 1$ to m_k , $m++$

- Compute ~~e_{mk}~~
$$e_{mk} = z_{mk} (1 - z_{mk}) \left(\sum_{i=1}^{m_{k+1}} \omega_{mi}^{(k)} e_i^{(k+1)} \right)$$

After repeating for each training sample

$$G_i(\omega) = \frac{\partial \mathcal{L}}{\partial \omega_{rs}^{(k)}} = \sum_{i=1}^n \frac{\partial \mathcal{L}_i}{\partial \omega_{rs}^{(k)}}$$

Using sum-of-squares error criterion,

$$G_i(w) = \frac{1}{2} [y_i - \hat{F}(x_i)]^2 \quad \text{for output } \vec{y} \text{ (different } h)$$

$$\therefore \frac{\partial G_i}{\partial a_0} = \hat{F}(x_i) - y_i$$

$$\frac{\partial G_i}{\partial a_m} = (\hat{F}(x_i) - y_i) B(x | \mu_m, \sigma_m)$$

$$\frac{\partial G_i}{\partial \mu_{jm}} = a_m (\hat{F}(x_i) - y_i) \frac{\partial B}{\partial \mu_{jm}}$$

$$= a_m (\hat{F}(x_i) - y_i) B(x | \mu_m, \sigma_m) (x_{ji} - \mu_{jm})$$

$$\frac{\partial G_i}{\partial \sigma_m} = \frac{a_m (\hat{F}(x_i) - y_i) B(x | \mu_m, \sigma_m) \sum_{j=1}^N (x_{ji} - \mu_{jm})^2}{\sigma_m^3}$$

$$\frac{\partial G_i}{\partial w} = \sum_{i=1}^N \frac{\partial G_i}{\partial w}$$

Here, we employed chain rule and each $\vec{x}_i = (x_{1i}, \dots, x_{ni}) \quad \forall i \in [1, N]$

Since Σ is a (symmetric) positive definite matrix, it has a Cholesky decomposition LL^T .

Let $y = x - \mu_m$, then

$$\begin{aligned} & (x - \mu_m)^T \Sigma (x - \mu_m) \\ &= y^T \Sigma y = y^T L L^T y \\ &= (L^T y)^T (L^T y) = Z^T Z \quad \therefore \text{Spherically Symmetric} \end{aligned}$$

$$\therefore \boxed{\tilde{x} = L^T x \text{ and } \tilde{\mu}_m = L^T \mu_m}$$

4. We claim that ~~Σ_m has~~ this property holds ~~for~~ if and only if Σ_m has exactly one eigenvalue λ_m with multiplicity 1.

\Leftarrow : Assuming Σ_m has exactly one eigenvalue λ_k with multiplicity one, then we can again write (using eigendecomposition) $\Sigma_m = S D S$:

$$\begin{aligned} & (x - \mu_m)^T \Sigma_m (x - \mu_m) \\ &= (x - \mu_m)^T S^T D S (x - \mu_m) \\ &= (S(x - \mu_m))^T \begin{bmatrix} 0 & & \\ & \lambda_m & \\ & & 0 \end{bmatrix} S(x - \mu_m) \end{aligned}$$

$\therefore B(x | \mu_m, \Sigma_m)$ only varies in direction of eigenvector corresponding to λ_k

\Rightarrow : Proceeding by contradiction, suppose the function varies in one direction only in the input space but ~~doesn't have~~ Σ_m has more than one eigenvalue. If this is the case, then the function can vary in any direction that is a linear combination of the eigenvectors corresponding to the eigenvalues, implying more than one direction, hence a contradiction. Therefore initial assumption is wrong and Σ_m has one eigenvalue with multiplicity 1. This completes the proof.

5

a.)

//Loading and labeling data.

```
> spam_all<-read.csv(file.choose())
> spam_test<-read.csv(file.choose())
> spam_train<-read.csv(file.choose())
> rflabs<-c("make","address","all","3d","our","over","remove","internet","order","mail","receive","
will","people","report","addresses","free","business","email","you","credit","your","front","000",
"money","hp","hpl","george","650","lab","labs","telnet","857","data","415","85","technology","19
99","parts","pm","direct","cs","meeting","original","project","re","edu","table","conference",";",
"{","[","!", "$", "#", "CAPAVE", "CAPMAX", "CAPTOT", "type")
> colnames(spam_all) = rflabs
> colnames(spam_test) = rflabs
> colnames(spam_train) = rflabs
```

//Standardizing predictor values in training data and isolating output.

```
> scaled_all = scale(spam_all)
> scaled_test = scale(spam_test)
> scaled_train = scale(spam_train)
> scaled_all = scale(spam_all)
> scaled_test = scale(spam_test)
> scaled_train = scale(spam_train)
> train_x = scaled_train[,c(1:57)]
> train_y = spam_train[,58]
```

//Running neural nets with number of edges being $59 \times (\text{size of hidden layer}) + 1$

```
> library(nnet)
```

```
> nnet_1 = nnet(train_x, train_y, size = 1, Wts = runif(60,-.5,.5), linout = T)
> nnet_2 = nnet(train_x, train_y, size = 1, Wts = runif(119,-.5,.5), linout = T)
> nnet_3 = nnet(train_x, train_y, size = 1, Wts = runif(178,-.5,.5), linout = T)
> nnet_4 = nnet(train_x, train_y, size = 1, Wts = runif(237,-.5,.5), linout = T)
> nnet_5 = nnet(train_x, train_y, size = 1, Wts = runif(296,-.5,.5), linout = T)
> nnet_6 = nnet(train_x, train_y, size = 1, Wts = runif(355,-.5,.5), linout = T)
> nnet_7 = nnet(train_x, train_y, size = 1, Wts = runif(414,-.5,.5), linout = T)
> nnet_8 = nnet(train_x, train_y, size = 1, Wts = runif(473,-.5,.5), linout = T)
> nnet_9 = nnet(train_x, train_y, size = 1, Wts = runif(532,-.5,.5), linout = T)
> nnet_10 = nnet(train_x, train_y, size = 1, Wts = runif(591,-.5,.5), linout = T)
```

Best model is, as expected, the one with 10 hidden units (lowest value after 100 iterations).

b.)

```
//Defining error matrix for all possible sizes of hidden layer and decay rate combinations.
```

```
err = matrix(rep(0,110), nrows = 10, ncols = 11)
```

```
//Standardizing predictor values in test data and isolating output.
```

```
test_x = scaled_test[,c(1:57)]
```

```
test_y = spam_test[,58]
```

```
//First loop is for units in hidden layer, the second is the decaying, and the third is the number of runs
```

```
for ( i in 1:10){
```

```
  for (j in 1:11){
```

```
    err_new = 0
```

```
    for (k in 1:10){
```

```
      nnet_l = nnet(train_x, train_y, size = i, Wts = runif(59*i+1,-.5,.5), linout = F, decay = 0.1*(j-1))
```

```
      y0 = predict(nnet_l, test_x, type = "raw")
```

```
//Converts into classified data
```

```
      y0 = ifelse(y0 < 0.5, 0, 1)
```

```
      err_new = err_new + mean(y0 != test_y)
```

```
    }
```

```
//Averaging error after runs and storing resulting value in matrix
```

```
      err[i,j]=err_new/10
```

```
    }
```

```
}
```

```
//Finding minimum error value in matrix. Row index corresponds to optimal size and column index corresponds to optimal regularization.
```

```
which (err == min(err),arr.ind = TRUE)
```

This results in 8 units and decay of 0.2 as optimal model.